

Adversarial Learning For End-To-End Cochlear Speech Denoising Using Lightweight Deep Learning Models

Tom Gajecki & Waldo Nogueira

Medical University Hannover, Cluster of Excellence ‘Hearing4all’, Hannover, Germany

ABSTRACT

This paper investigates an end-to-end speech signal denoising approach for cochlear implants (CIs). Building on previous work, we first explore the effect of relocating the deep envelope detector within the deep learning-based CI sound coding strategy, moving it from the skip connection to the output of the masking operation. This modification enables high-resolution time-frequency masking and optimizes noise reduction. Next, we introduce a discriminator network to further enhance the model by enforcing the generation of higher-quality electrograms (i.e., the electric pulse patterns that stimulate the auditory nerve). This adversarial learning approach improves the generation of electrograms and has the potential to enhance speech understanding for CI users. Objective evaluations, including signal-to-noise ratio improvement and linear cross-correlation coefficients, demonstrate that these enhancements significantly boost the performance of the end-to-end CI speech-denoising algorithm while reducing its parameter count, making it suitable for real-time applications.

Index Terms— Cochlear Implants, Deep Neural Networks, Adversarial Training, End-To-End, Speech Enhancement

1. INTRODUCTION

Cochlear implants (CIs) have long been a solution for individuals with profound sensorineural hearing loss, enabling them to regain the sensation of hearing through direct electrical stimulation of the auditory nerve [1]. Despite significant advancements in CI technology, a persistent challenge remains in noisy environments, where traditional sound coding strategies, such as the Advanced Combination Encoder (ACE), often struggle to maintain speech intelligibility. Recent approaches utilizing deep learning for end-to-end speech denoising in CIs have shown great potential in enhancing noise reduction by generating high-quality electrograms directly from raw audio signals [2, 3]. However, many of these models are computationally intensive, which can hinder their real-time applicability.

The funding for this work was provided by the German Research Foundation (DFG) under Project ID 381895691, led by Waldo Nogueira.

In this paper, we present an extension of an end-to-end CI speech denoising sound coding strategy [4]. While the end-to-end CI processing explored here can be applied to any commercial CI sound coding strategy, our focus is on ACE (i.e. Deep ACE). We begin by examining the impact of relocating the deep envelope detector (DED) from the skip connection to the output of the masking operation. Next, we introduce an adversarial version of our end-to-end speech denoising system, following a similar approaches such as [5, 6]. This new model contains a discriminator and acts as an adversarial network, helping to ensure that the electrograms are both noise-reduced and better aligned with the underlying speech signals.

Furthermore, we provide a theoretical analysis suggesting that electrograms are inherently less complex than raw audio signals. This reduction in complexity makes electrograms easier for neural networks to predict, as the network can focus on capturing the most relevant features for speech intelligibility without being overwhelmed by the intricacies of raw audio. By simplifying the prediction task, the network can allocate more resources to refining the electrogram quality, thereby improving its ability to preserve crucial speech information.

The proposed model is evaluated through a series of objective measures, demonstrating that the adversarial deep learning-based CI speech denoising model outperforms the used baselines. This approach offers a more efficient solution for potential real-time deployment, with improvements observed in noise reduction and the quality of generated electrograms. Although listening tests were not conducted, the objective results suggest that our model could significantly enhance speech intelligibility and overall user experience in complex auditory environments.

2. METHODS & MATERIALS

2.1. Algorithms

Deep ACE V2: This is an enhanced version of the original Deep ACE model [4] (see Figure 1a), designed to improve noise reduction and electrogram quality for CIs. A key innovation in Deep ACE V2 is the repositioning of the Deep Envelope Detector (DED; Figure 1b) to occur after the mask-

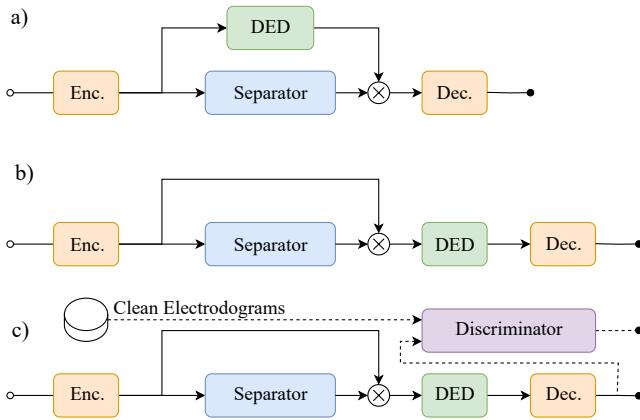


Fig. 1. Deep learning modules of the investigated end-to-end CI speech denoising variants. The original Deep ACE with the deep envelope detector (DED) in the skip connection (a), the here presented Deep ACE V2 with the DED after the masking operation (b), and the Adversarial Deep ACE (c).

ing operation. This adjustment allows the model to leverage high-resolution time-frequency masking, capturing more detailed speech features and filtering out noise more effectively. By placing the DED after the masking stage, Deep ACE V2 generates electrograms that are more refined and closely aligned with the desired outputs, leading to better speech intelligibility in noisy environments. This configuration also lays the foundation for further enhancements, such as the integration of a discriminator network, as explored in this study.

Adversarial Deep ACE model: In this model, a multi-scale convolutional discriminator network (inspired by [7]) is introduced to enhance electrogram generation (see Figure 1c). Acting as an adversarial network, the discriminator distinguishes between realistic and unrealistic electrograms, guiding the generator to produce higher-quality outputs that better mimic auditory signals. The discriminator operates at multiple resolutions, using downsampling to capture both fine and coarse details while applying normalization techniques to ensure stable training. Notably, the discriminator is only used during training, maintaining a low computational footprint for real-time inference. Combined with the strategic placement of the DED after the masking operation, this approach enables Adversarial Deep ACE to leverage high-resolution time-frequency masking, resulting in improved noise reduction and enhanced speech intelligibility in noisy environments.

2.2. Loss Functions for adversarial Speech Denoising

The training of the Adversarial Deep ACE model involves several key loss functions that collectively contribute to generating high-quality electrograms for CIs. The training

configuration for the other models used in this study is consistent with the procedures described in [4].

Adversarial Loss (\mathcal{L}_{adv}): The adversarial loss is crucial for training the generator to produce output signals that are indistinguishable from real ones by the discriminator. It is calculated using binary cross-entropy (BCE), averaged over all discriminators. Specifically, the adversarial loss is defined as:

$$\mathcal{L}_{adv} = \frac{1}{N} \sum_{i=1}^N \text{BCE}(\mathbf{1}, D_i(G(x))), \quad (1)$$

where N represents the number of discriminators, $D_i(G(x)_i)$ denotes the discriminator's output for the generated denoised speech from the noisy audio signal x at the i -th discriminator, and $\mathbf{1}$ is a vector of ones, representing the target label for real target data.

Reconstruction Loss (\mathcal{L}_{recon}): The reconstruction loss measures the discrepancy between the target signals and those generated by the model. In Adversarial Deep ACE, this loss is computed using the Mean Squared Error (MSE) between the generated $G(x)$ and the target signals. This ensures that the generated signals closely align with the target signals, improving the quality of the output electrograms. In contrast, Adversarial TasNet uses the Scale-Invariant Signal-to-Distortion Ratio (SI-SDR; [8]) as its reconstruction loss, optimizing the generated signals to better match the clean target audio.

Feature Matching Loss (\mathcal{L}_{feat_match}): Feature matching loss [9] is introduced to encourage the generator to produce electrograms that match the feature space of real electrograms. This is calculated as:

$$\mathcal{L}_{feat_match} = \frac{1}{N_D} \sum_{i=1}^{N_D} \frac{4}{n_{\text{layers}} + 1} \sum_{j=1}^{n_{\text{layers}}} \left\| D_{\text{real}}^{(i,j)} - D_{\text{fake}}^{(i,j)} \right\|_1 \quad (2)$$

Here, N_D is the number of discriminators, and n_{layers} is the number of layers within each discriminator. The outputs $D_{\text{real}}^{(i,j)}$ and $D_{\text{fake}}^{(i,j)}$ represent the discriminator's response to real and generated electrograms at the i -th discriminator and j -th layer, respectively. The L_1 norm is used to measure the absolute difference between these features, ensuring that the generated electrograms not only look realistic but also possess similar internal feature structures as real electrograms.

Total Loss: The total loss combines adversarial, reconstruction, and feature-matching losses, weighted empirically as 10.0 for feature matching and 1.0 for the others, which balance contributions to ensure high-quality outputs.

The models in this work are designed to be lightweight, as summarized in Table 1. For a proper baseline comparison with Adversarial Deep ACE, we employ an adversarial version of TasNet, which uses the same discriminator as the one in Adversarial Deep ACE and is trained similarly. The table provides the number of parameters for both the generator and, where applicable, the discriminator. It is important to highlight that the discriminator is only utilized during training to improve the quality of the generated electrograms, but it is not required during inference. This contributes to the system’s efficiency. By maintaining a low parameter count, these models strike a balance between computational efficiency and performance, making them well-suited for real-time applications.

	#Params gen.	#Params disc.
TasNet [10]	164,736	-
Adv. TasNet	164,736	97,575
Deep ACE [4]	172,934	-
Deep ACE V2	172,934	-
Adv. Deep ACE	172,934	189,660

Table 1. Comparison of Generator (gen.) and Discriminator (disc.) parameter counts for various models. This table presents the number of parameters for both the generator and discriminator across different models.

2.3. Objective Instrumental Measures

SNRi: To evaluate the noise reduction achieved by each of the tested algorithms, we compute the SNR improvement (SNRi). This metric is calculated in the electrogram domain and compares the original input SNR with the SNR obtained after denoising. The SNRi is defined as:

$$\text{SNRi} = 10 \cdot \log_{10} \left(\frac{\sum_{k=1}^M \|\mathbf{p}_k^n - \mathbf{p}_k^c\|^2}{\sum_{k=1}^M \|\mathbf{p}_k^d - \mathbf{p}_k^c\|^2} \right), \quad (3)$$

where \mathbf{p}_k represents the LGF output of band k , and the superscripts n , c , and d correspond to the noisy, clean, and denoised electrograms, respectively.

Similarly, SNRi can also be computed in the audio domain by replacing the LGF output with the predicted audio signals. This involves comparing the SNR of the original noisy audio signal to the SNR of the denoised audio output, using the predicted audio in place of electrograms.

Linear Cross Correlation (LCC): To evaluate distortions and artifacts from the tested algorithms, we calculated the linear correlation coefficients (LCCs) between the clean ACE electrograms (\mathbf{p}^c) and the denoised electrograms (\mathbf{p}^d). The LCCs were computed for each of the 22 channels to assess degradation in individual channel outputs caused by the denoising process. The LCC for channel k (LCC_k) is

calculated using the Pearson correlation coefficient [11] as follows:

$$\text{LCC}(p_k, p_{k,\text{ref}}) = \frac{\text{Cov}(p_k, p_{k,\text{ref}})}{\sigma_{p_k} \sigma_{p_{k,\text{ref}}}} \quad (4)$$

where $\text{Cov}(p_k, p_{k,\text{ref}})$ is the covariance between the predicted and reference electrograms at frame k , and σ_{p_k} and $\sigma_{p_{k,\text{ref}}}$ are their standard deviations. A higher LCC indicates a stronger linear relationship, reflecting the predicted electrograms’ accuracy in capturing the reference signals’ underlying patterns. For the TasNet variants, the LGF output is replaced by the estimated denoised speech audio signal.

Complexity Measures: To assess the complexity of the predicted electrograms and audio samples, we employ two complementary measures: *Lempel-Ziv complexity* and *Shannon entropy*. Each measure provides a different perspective on the variability and unpredictability of the sequences.

Lempel-Ziv complexity (LZC) quantifies the complexity of the predicted electrogram sequences and audio samples by counting the number of distinct patterns or configurations observed throughout the sequence. For electrograms, LZC measures the diversity of electrode activation patterns, while for audio signals, it evaluates the occurrence of new, previously unseen patterns in the time-domain waveform. A higher LZC value indicates a more complex and less predictable sequence, reflecting greater variability and richness in both the electrograms and audio signals.

Shannon entropy measures the unpredictability or information content in the predicted electrogram sequences and audio samples. Entropy is calculated as:

$$H = - \sum_i q_i \log_2(q_i) \quad (5)$$

where q_i represents the probability of occurrence for each unique pattern. In electrograms, entropy assesses the distribution of electrode activation patterns across frames, while in audio samples, it evaluates the variability in signal amplitude or frequency content. Higher entropy values suggest more complex and less predictable sequences, indicating greater variability in both electrograms and audio signals.

2.4. Audio Data

In this study, we trained the models using speech signals from the LibriVox corpus [12] and mixed them with noise from the DEMAND dataset [13]. For testing, we utilized the Hochmair, Schulz, Moser (HSM) sentence dataset, which was mixed with CCIT speech-shaped noise [14] and ICRA7 babble noise [15]. The noise levels were adjusted to a signal-to-noise ratio (SNR) range of -5 to 10 dB.

3. RESULTS

SNR_i: Figure 2 presents the SNR_i achieved by the evaluated algorithms. The results show that the proposed Adversarial Deep ACE method consistently outperforms the baseline algorithms across all input SNRs for both CCITT and ICRA7 noise conditions. Interestingly, the adversarial TasNet performs worse than the standard TasNet in CCITT noise but performs comparably in ICRA7. This discrepancy may be attributed to suboptimal tuning of the adversarial model for specific noise types, indicating the need for further optimization or task-specific adjustments.

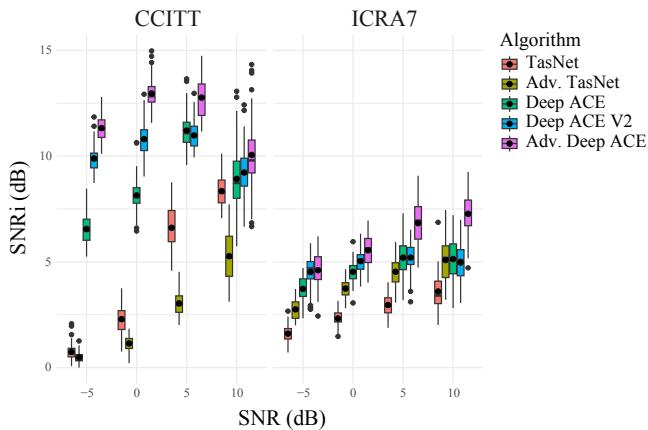


Fig. 2. Box plots showing the SNR_i scores in dB for the tested algorithms in SSN and ICRA7 noises for the different SNRs using the HSM speech dataset.

LCCs: Figure 3 shows the LCCs as a function of electrode number for each algorithm. The results indicate that CI denoising strategies perform best at lower frequencies, likely due to their focus on enhancing low-frequency modulations, which are critical for speech understanding [16]. In contrast, TasNet shows a slight advantage at higher frequencies, which are important for capturing consonants and environmental sounds.

Complexity measures: Figure 4 shows violin plots comparing the complexity of audio signals and electrograms. The complexity measures, based on LZC and Shannon entropy, reveal that electrograms have lower complexity than audio signals. This simplification likely explains the superior performance of CI denoising strategies, which focus on predicting electrograms. In contrast, TasNet, which directly predicts audio, must handle higher complexity, making denoising more challenging in noisy environments. Predicting electrograms offers a more manageable signal, potentially improving CI denoising effectiveness.

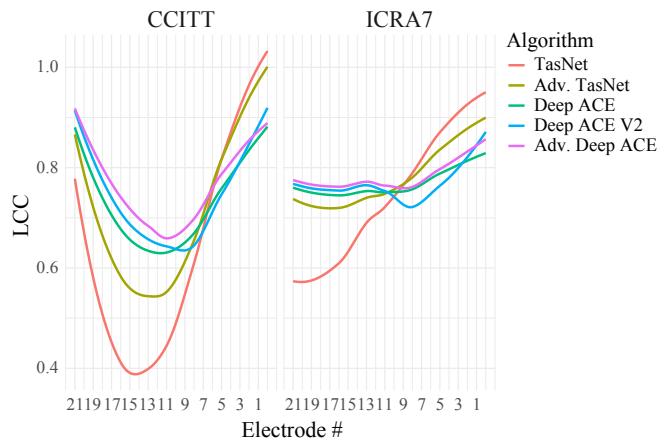


Fig. 3. Polynomial regressions showing the channel-wise LCCs between processed and clean electrograms for the different algorithms and noises using the HSM dataset. Higher electrode numbers correspond to higher frequencies.




Fig. 4. Violin plots comparing the complexity of audio signals and electrograms using Lempel-Ziv complexity (LZC) and Shannon entropy.

4. CONCLUSIONS

In this study, we enhanced a deep learning-based denoising CI sound coding strategy by integrating a discriminator network and repositioning the DED, improving noise reduction and electrogram quality. Objective measures like SNR improvement and LCCs show that the adversarial CI sound coding strategy outperforms the used baselines. The simplicity of electrograms compared to raw audio likely contributes to this improved performance. These advancements represent a promising step toward more effective, real-time speech enhancement in noisy environments. Future work could focus on the optimization and validation through listening tests.

5. REFERENCES

- [1] Thomas Lenarz, “Cochlear implant—state of the art,” *Laryngorhinootologie*, vol. 96, no. 1, pp. 123–151, 2017.
- [2] Tom Gajecki and Waldo Nogueira, “An end-to-end deep learning speech coding and denoising strategy for cochlear implants,” in *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2022, pp. 3109–3113.
- [3] Enoch Hsin-Ho Huang, Rong Chao, Yu Tsao, and Chao-Min Wu, “Electrodenet—a deep-learning-based sound coding strategy for cochlear implants,” *IEEE Transactions on Cognitive and Developmental Systems*, vol. 16, no. 1, pp. 346–357, 2024.
- [4] Tom Gajecki, Yichi Zhang, and Waldo Nogueira, “A deep denoising sound coding strategy for cochlear implants,” *IEEE Transactions on Biomedical Engineering*, vol. 70, no. 9, pp. 2700–2709, 2023.
- [5] Chengyun Deng, Yi Zhang, Shiqian Ma, Yongtao Sha, Hui Song, and Xiangang Li, “Conv-tassan: Separative adversarial network based on conv-tasnet.,” in *Interspeech*, 2020, pp. 2647–2651.
- [6] Santiago Pascual, Antonio Bonafonte, and Joan Serra, “Segan: Speech enhancement generative adversarial network,” *arXiv preprint arXiv:1703.09452*, 2017.
- [7] Kundan Kumar, Rithesh Kumar, Thibault De Boissiere, Lucas Gestin, Wei Zhen Teoh, Jose Sotelo, Alexandre De Brebisson, Yoshua Bengio, and Aaron C Courville, “Melgan: Generative adversarial networks for conditional waveform synthesis,” *Advances in neural information processing systems*, vol. 32, 2019.
- [8] J. L. Roux, S. Wisdom, H. Erdogan, and J. R. Hershey, “SDR – half-baked or well done?,” in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2019, pp. 626–630.
- [9] Tim Salimans, Ian Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen, “Improved techniques for training gans,” *Advances in neural information processing systems*, vol. 29, 2016.
- [10] Yi Luo and Nima Mesgarani, “Conv-TasNet: Surpassing ideal time-frequency magnitude masking for speech separation,” *IEEE/ACM transactions on audio, speech, and language processing*, vol. 27, no. 8, pp. 1256–1266, 2019.
- [11] David Freedman, Robert Pisani, and Roger Purves, “Statistics,” *Pisani, R. Purves, 4th edn. WW Norton & Company, New York*, 2007.
- [12] Benjamin Beilharz, Xin Sun, Sariya Karimova, and Stefan Riezler, “LibriVoxDeEn: A corpus for german-to-english speech translation and speech recognition,” in *Proceedings of LREC*, 2020.
- [13] Joachim Thiemann, Nobutaka Ito, and Emmanuel Vincent, “DEMAND: a collection of multi-channel recordings of acoustic noise in diverse environments,” in *Proc. Meetings Acoust*, 2013, pp. 1–6.
- [14] Hugo Fastl and Eberhard Zwicker, “Psychoacoustics - facts and models,” *Springer*, vol. 3rd edition, 2007.
- [15] Wouter A Dreschler, Hans Verschuure, Carl Ludvigsen, and Søren Westermann, “ICRA noises: Artificial noise signals with speech-like spectral and temporal properties for hearing instrument assessment,” *Audiology*, vol. 40, no. 3, pp. 148–157, 2001.
- [16] Bas van Dijk, Marc Moonen, Jan Wouters, et al., “Understanding the effect of noise on electrical stimulation sequences in cochlear implants and its impact on speech intelligibility,” *Hearing Research*, vol. 299, pp. 79–87, 2013.